

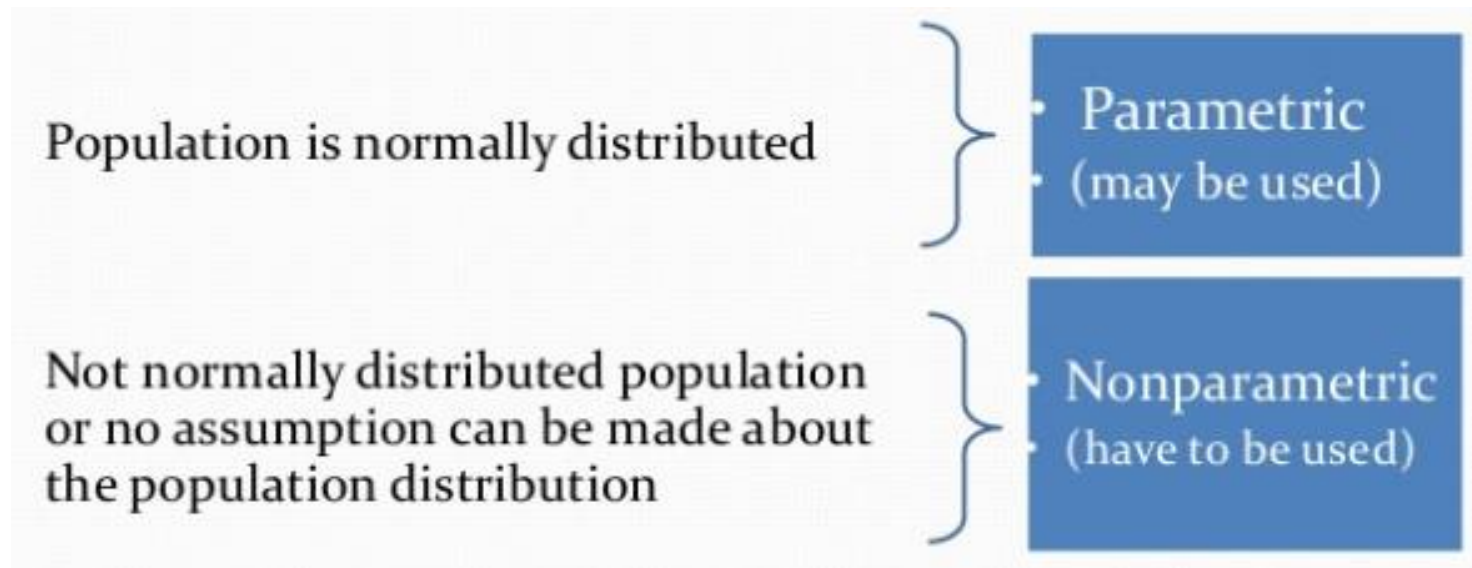
Pattern Recognition

Krystian Mikolajczyk

Blackboard

Data representation

Parametric vs non parametric



Parametric vs non parametric

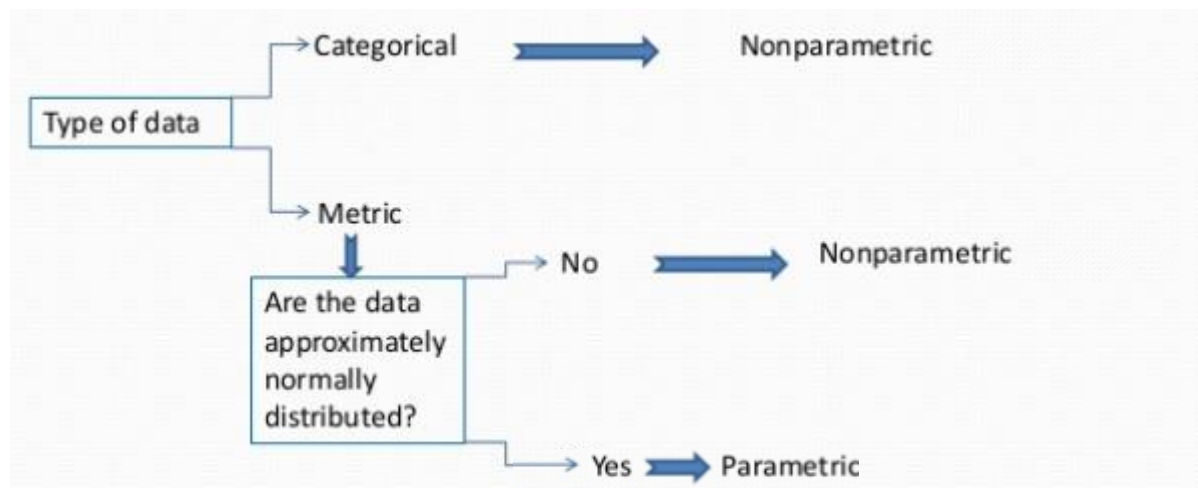
Normal Distribution

- a very common continuous probability distribution
- All normal distributions are symmetric.
- bell-shaped curve with a single peak.
- **68%** of the observations fall within **1 standard deviation** of the **mean**
- **95%** of the observations fall within **2 standard deviations** of the **mean**
- **99.7%** of the observations fall within **3 standard deviations** of the **mean**
- for a normal distribution, almost all values lie within **3 standard deviations** of the mean

Parametric vs non parametric

- In cases where
 - the data which are measured by interval or ratio scale come from a normal distribution *
parametric are used.
- In cases where
 - the data is nominal or ordinal
 - the assumptions of parametric tests are inappropriate
nonparametric are used.

* - simplifying assumption,
other distributions are possible



Parametric




- Multivariate Gaussian

$$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|^{1/2}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}$$

mean covariance

$$\boldsymbol{\mu}_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

$$\boldsymbol{\Sigma}_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \boldsymbol{\mu}_{\text{ML}})(\mathbf{x}_n - \boldsymbol{\mu}_{\text{ML}})^\top$$



Parametric – Gaussian mixture

- Linear super-position of Gaussians

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

- Normalization and positivity require

$$\sum_{k=1}^K \pi_k = 1 \quad 0 \leq \pi_k \leq 1$$




- Can interpret the mixing coefficients as prior probabilities

$$p(\mathbf{x}) = \sum_{k=1}^K p(k) p(\mathbf{x} | k)$$


Parametric – Gaussian mixture

- 3 Gaussians

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$



Fisher Vectors



FV encoding $\phi =$

+ sqrt-l²
normalisation

$$\begin{bmatrix} \mathbf{v}_1 \\ \mathbf{u}_1 \\ \mathbf{v}_2 \\ \mathbf{u}_2 \\ \vdots \\ \mathbf{v}_K \\ \mathbf{u}_K \end{bmatrix}$$

first and second order statistics


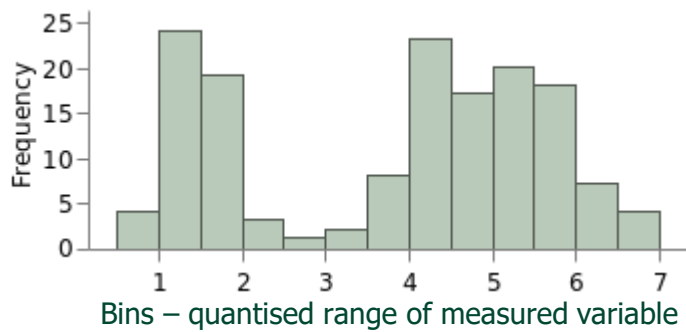
$$\mathbf{v}_k = \frac{1}{M\sqrt{\pi_k}} \sum_{i=1}^M \gamma_k(\mathbf{x}_i) \frac{\mathbf{x}_i - \mu_k}{\sigma_i}$$

$$\mathbf{u}_k = \frac{1}{M\sqrt{2\pi_k}} \sum_{i=1}^M \gamma_k(\mathbf{x}_i) \left(\frac{\mathbf{x}_i - \mu_k}{\sigma_i} - 1 \right)^2$$

Nonparametric - histogram

- A histogram is a representation of the distribution of data.
 - an estimate of the probability distribution of a continuous variable
 - a rough sense of the density of the underlying distribution of the data
 - total area of a histogram used for probability density is always normalized to 1
 - m_i - bin value
 - n - number of observations
 - k - number of bins



$$n = \sum_{i=1}^k m_i.$$



Nonparametric - histogram

- Cumulative histogram - counts the cumulative number of observations in all of the bins up to the specified bin

$$M_i = \sum_{j=1}^i m_j$$



Nonparametric - histogram

• Quantisation of data space

- k – number of bins
- x – data point
- h – bin width

$$k = \left\lceil \frac{\max x - \min x}{h} \right\rceil$$

general

$$k = \sqrt{n}$$



assumes normal distribution

$$k = \lceil \log_2 n \rceil + 1$$

$$k = \lceil 2n^{1/3} \rceil$$

non-normal data + adds bins for skewness

$$k = 1 + \log_2(n) + \log_2 \left(1 + \frac{|g_1|}{\sigma_{g_1}} \right) \text{Skew}$$



$$\text{Skew} = g_1 = \frac{\sum (x_i - \mu)^3}{n\sigma^3}$$


$$\sigma_{g_1} = \sqrt{\frac{6(n-2)}{(n+1)(n+3)}} \quad 2$$

Clustering


Clustering

- Classification of similar objects into different groups
- partitioning of a data set into subset (clusters), so that the data in each subset (ideally) share some common trait - often proximity according to some defined distance measure.

Clustering



Clustering



Clustering

- kmeans - partitional
 - Top-down
- Agglomerative
 - Bottom-up


Top-down clustering

- Kmeans

- Random initialization with arbitrarily set K - initial cluster centers

For each data point find the nearest cluster center and assign the point to this cluster

Recompute the cluster center by computing the average of the







Top-down clustering

- Kmeans

- Random initialization with arbitrarily set K



The solution depends on initialization – local optimum



Top-down clustering

- Kmeans


- Random initialization with arbitrarily set K
- finds local optimum
- sensitive to outliers
- cluster centers might not represent well the point distribution
- $O(Nkld)$ - low complexity, efficient
 - Not if k is comparable with N
- Good method for partitioning the data



Bottom-up agglomerative

Set the threshold distance – maximum distance between two points


1. Compare all pairs and mark the with the smallest distance
2. Compare with the threshold – stop clustering if distance is larger than the threshold
3. Merge the pair and replace by its average
4. Repeat points 1 to 3



Bottom-up agglomerative

- Agglomerative,

- Arbitrarily set maximum cluster size
- Robust to outliers
- Meaningful cluster centers



Distance/Similarity Metrics

Intro

- Numerical data

- Each data instance is a numerical feature vector.


$$\mathbf{x} = \begin{pmatrix} 26 \\ 21.6 \\ 102 \\ \dots \end{pmatrix}$$

- Example: the age, body mass index, blood pressure, ... of a patient.

- Each instance is a structured object: a string, a tree or a graph.

- Examples: words, DNA sequences, XML documents, molecules, social communities...

SFGAAJHSKJH



Intro


- Pairwise metric
 - Informally, a way of measuring the distance (or similarity) between object
- Metrics are ubiquitous in machine learning
 - Get yourself a good metric and you've basically solved the problem. Metrics are convenient proxies to manipulate complex objects.

Applications

- Classification
- k-Nearest Neighbors,
- Support Vector Machines...
- Clustering: K-Means and its variants.
- Information Retrieval / Ranking: search by query, image/document retrieval...
- Data visualization in high dimensions. ...
- Computer Vision: compare images or videos in ad-hoc representations. Used in image classification, object/face recognition, tracking, image annotation...
- Bioinformatics: compare structured objects such as DNA sequences or temporal series.
- Whenever the notion of metric plays an important role. Fun examples of applications include music recommendation, identity verification, cartoon synthesis and assessing the efficacy of acupuncture, to name a few

Metrics

- The notion of good metric is problem-dependent
 - Each problem has its own semantic notion of similarity, which is often badly captured by standard metrics (e.g., Euclidean distance).
- Solution: learn the metric from data
 - Basic idea: learn a metric that assigns small (resp. large) distance to pairs of examples that are semantically similar (resp. dissimilar).



Metric

- Non-negativity:
$$D(P, Q) \geq 0$$
- Identity of indiscernibles:
$$D(P, Q) = 0 \text{ iff } P = Q$$
- Symmetry:
$$D(P, Q) = D(Q, P)$$
- Subadditivity (triangle inequality):
$$D(P, Q) \leq D(P, K) + D(K, Q)$$

Pseudo-Metric (aka Semi-Metric)

- Non-negativity:
$$D(P, Q) \geq 0$$
- Property changed to:
$$D(P, Q) = 0 \text{ if } P = Q$$
- Symmetry:
$$D(P, Q) = D(Q, P)$$
- Subadditivity (triangle inequality):
$$D(P, Q) \leq D(P, K) + D(K, Q)$$

Similarity measures

- String edit distance

- The edit distance is the cost of the cheapest sequence of operations (script) turning a string into another. Allowable operations are insertion, deletion and substitution of symbols. Costs are gathered in a matrix C .

C		\$	a	b
\$		0	1	1
a		1	0	1
b		1	1	0

\Rightarrow edit distance between abb and aa is 2 (needs at least two operations)

- It is a proper distance if and only if C satisfies:

$$C_{ij} \geq 0, \quad C_{ij} = C_{ji}, \quad C_{ik} \leq C_{ij} + C_{jk} \quad \forall i, j, k$$

- Generalization to trees (quadratic or cubic complexity) and graphs (NP-complete).

Minkowski-Form Distances

$$L_p(P, Q) = \left(\sum_i |P_i - Q_i|^p \right)^{\frac{1}{p}}$$

$$L_1(P, Q) = \sum_i |P_i - Q_i|$$


$$L_2(P, Q) = \sqrt{\sum_i (P_i - Q_i)^2}$$

$$L_\infty(P, Q) = \max_i |P_i - Q_i|$$

Similarity Measures

- Euclidean Distance

$$d(Q, V) = \sum_i (q_i - v_i)^2$$



- Properties

- Focuses on the differences between the histograms
- Range: $[0, \infty]$
- All cells are weighted equally.
- Not very robust to outliers!

Similarity measures

- Cosine similarity

- The cosine similarity measures the cosine of the angle between two instances, and can be computed a

$$K_{\cos}(\mathbf{x}, \mathbf{x}') = \frac{\mathbf{x}^T \mathbf{x}'}{\|\mathbf{x}\|_2 \|\mathbf{x}'\|_2}.$$

- It is widely used in data mining (better notion of similarity for bag-of-words + efficiently computable for sparse vectors).

- Bilinear similarity

- The bilinear similarity is related to the cosine but does not include normalization and is parameterized by a matrix \mathbf{M} :

$$K_{\mathbf{M}}(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{M} \mathbf{x}',$$


- where $\mathbf{M} \in \mathbb{R}^{d \times d}$ is not required to be PSD nor symmetric

Similarity Measures

- Cross correlation



$$f_1 = \begin{bmatrix} I_1(0,0) \\ I_1(0,1) \\ I_1(0,2) \\ \vdots \end{bmatrix}$$



$$f_2 = \begin{bmatrix} I_2(0,0) \\ I_2(0,1) \\ I_2(0,2) \\ \vdots \end{bmatrix}$$

$$D_{Euclidean}^2(f_1, f_2) = \|f_1 - f_2\| = \sum_n (f_1(n) - f_2(n))^2 =$$

If the patches are normalized then $\sum_n f_1^2(n) = \sum_n f_2^2(n) = \text{constant}$

$$\text{CrossCorrelation}(f_1, f_2) = \sum_n f_1(n) f_2(n)$$

- Properties


- very distinctive,
- easy to implement

- very high dimensional,
- not very robust to noise

Similarity Measures

- Intersection

$$\cap(Q, V) = \sum_i \min(q_i, v_i)$$



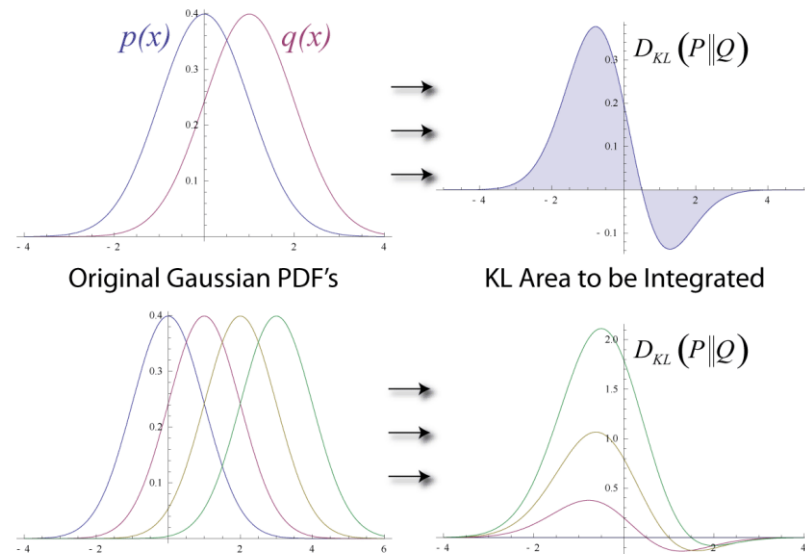
- Properties

- Measures the common part of both histograms
- Range: $[0,1]$
- For unnormalized histograms, use the following formula

$$\cap(Q, V) = \frac{1}{2} \left(\frac{\sum_i \min(q_i, v_i)}{\sum_i q_i} + \frac{\sum_i \min(q_i, v_i)}{\sum_i v_i} \right)$$

Kullback-Leibler Divergence

$$KL(P, Q) = \sum_i P_i \log \frac{P_i}{Q_i}$$



- Properties

- Difference between two probability distributions
- Information theoretic origin
- Non symmetric
- Not a metric
- $Q_i = 0$?

Jensen-Shannon Divergence

$$JS(P, Q) = \frac{1}{2}KL(P, M) + \frac{1}{2}KL(Q, M)$$


$$M = \frac{1}{2}(P + Q)$$

$$JS(P, Q) = \frac{1}{2} \sum_i P_i \log \frac{2P_i}{P_i + Q_i} +$$

$$\frac{1}{2} \sum_i Q_i \log \frac{2Q_i}{P_i + Q_i}$$

Properties

- Information theoretic origin
- Symmetric.
- \sqrt{JS} is a metric.



Jensen-Shannon Divergence

- Using Taylor extension and some algebra:


$$\begin{aligned} JS(P, Q) &= \sum_{n=1}^{\infty} \frac{1}{2n(2n-1)} \sum_i \frac{(P_i - Q_i)^{2n}}{(P_i + Q_i)^{2n-1}} \\ &= \frac{1}{2} \sum_i \frac{(P_i - Q_i)^2}{(P_i + Q_i)} + \\ &\quad \frac{1}{12} \sum_i \frac{(P_i - Q_i)^4}{(P_i + Q_i)^3} + \dots \end{aligned}$$

$$\chi^2(P, Q) = \frac{1}{2} \sum_i \frac{(P_i - Q_i)^2}{(P_i + Q_i)}$$

Similarity Measures

- Chi-square

$$\chi^2(P, Q) = \frac{1}{2} \sum_i \frac{(P_i - Q_i)^2}{(P_i + Q_i)}$$



- Properties


- Statistical background:
 - Test if two distributions are different
 - Experimentally results are very similar to JS.
- Range: $[0, \infty]$
- Cells are not weighted equally!
- $\sqrt{\chi^2}$ is a metric, experimentally better than L_2
- More robust to outliers, reduces the effect of large bins

χ^2 Histogram Distance

- Different metrics may lead to opposite conclusion




- Bin-to-Bin distances such as L2 are sensitive to quantization:



Bin-to-Bin Distances

- #bins  \Rightarrow robustness  distinctiveness 
- #bins  \Rightarrow robustness  distinctiveness 



Can we achieve **robustness**
and
distinctiveness ?

Quadratic-Form Histogram Distance

$$\begin{aligned} QF^A(P, Q) &= \sqrt{(P - Q)^T A (P - Q)} \\ &= \sqrt{\sum_{ij} (P_i - Q_i)(P_j - Q_j) A_{ij}} \end{aligned}$$

- Properties

- A_{ij} is the similarity between bin i and j .
- If A is the inverse of the covariance matrix, QF is called Mahalanobis distance.
- If A is the identity matrix, QF is the Euclidean distance

$$A = I \rightarrow \sqrt{\sum_{ij} (P_i - Q_i)^2} = \sqrt{L_2(P, Q)}$$

Quadratic-Form Histogram Distance

$$QFA(P, Q) = \sqrt{(P - Q)^T A (P - Q)}$$

- Properties
 - Does not reduce the effect of large bins.
 - Alleviates the quantization problem.
 - Linear time computation in # of non zero A_{ij}
 - If A is positive-definite then QF is a metric
 - If A is positive-semidefinite then QF is a pseudo-metric

Quadratic-Form Histogram Distance

$$\begin{aligned} QFA(P, Q) &= \sqrt{(P - Q)^T A (P - Q)} \\ &= \sqrt{(P - Q)^T W^T W (P - Q)} \\ &= L_2(WP, WQ) \end{aligned}$$

- Properties

- We assume there is a **linear transformation** that makes bins **independent**
 - There are cases where this is **not true** *e.g.* **COLOR**
- Converting distance to similarity (Hafner *et. al* 95):

$$A_{ij} = 1 - \frac{D_{ij}}{\max_{ij}(D_{ij})} \quad \text{or}$$

$$A_{ij} = e^{-\alpha \frac{D(i,j)}{\max_{ij}(D(i,j))}}$$

If α is large enough,
 A will be positive-definitive


Quadratic-Chi Histogram Distance

$$QC_m^A(P, Q) = \sqrt{\sum_{ij} \frac{(P_i - Q_i)(P_j - Q_j)A_{ij}}{(\sum_c (P_c + Q_c)A_{ci})^m (\sum_c (P_c + Q_c)A_{cj})^m}}$$


Properties

- A_{ij} is the similarity between bin i and j .
- Generalizes QF and χ^2
- Reduces the effect of large bins.
- Alleviates the quantization problem.
- Linear time computation in # non zero A_{ij} .
- non-negative if A is positive-semidefinite.
- Symmetric
- Triangle inequality unknown.
- If we define $\frac{0}{0} = 0$ and $0 \leq m < 1$, QC is continuous.

Similarity-Invariance







Sparseness-Invariance



The Earth Mover's Distance

The Earth Mover's Distance


- The Earth Mover's Distance is defined as the minimal cost that must be paid to transform one histogram into the other, where there is a “ground distance” between the basic features that are aggregated into the histogram.

 \neq  $=$ 

Earth Mover's Distance - example

[Rubner *et al.* 1998]

- We are given the distribution of grades for a course over the past three years and we want to compare the distributions:



- If we just compare theses as vectors, the results from Year 3 are as similar to the results from Year 2 as they are to the results of Year 1.

Earth Mover's Distance - approach

- Instead of comparing the values in each bin (e.g. Euclidean distance), compute the amount of work needed to transform one distribution into the other.
- Define the cost of moving d values from bin i to bin j as:

$$\text{Work}(d, i, j) = d \cdot |i - j|$$

- Find the minimal amount of work that needs to be done to transform one distribution into the other.

Earth Mover's Distance - definition

- Given distributions $X=\{x_1,\dots,x_n\}$ and $Y=\{y_1,\dots,y_n\}$, set $c_{ij}=|i-j|$ and find the values for f_{ij} that minimize:

$$\text{EMD}_{\text{work}} = \min_{F=\{F_{ij}\}} \sum_{i=1}^n \sum_{j=1}^n f_{ij} c_{ij}$$

- subject to the constraints:


$c_{ij} \geq 0$ cost of moving
from bin i to bin j

$f_{ij} \geq 0$ amount of data moved
from bin i to bin j

$$\sum_{j=1}^n f_{ij} = x_i \quad \text{and} \quad \sum_{i=1}^n f_{ij} = y_j$$

Earth Mover's Distance - illustration


- In general, this is the *transportation problem* and can be solved using linear programming.
- For 1D histograms, this can be solved using the greedy algorithm.



Earth Mover's Distance - illustration

- In general, this is the *transportation problem* and can be solved using linear programming.
- For 1D histograms, this can be solved using the greedy algorithm.


- Find the first non-empty bin.
- Move earth into first non-empty bin in the other histogram.



Earth Mover's Distance - illustration

- In general, this is the *transportation problem* and can be solved using linear programming.
- For 1D histograms, this can be solved using the greedy algorithm.


- Find the first non-empty bin.
- Move earth into first non-empty bin in the other histogram.



Earth Mover's Distance - illustration

- In general, this is the *transportation problem* and can be solved using linear programming.
- For 1D histograms, this can be solved using the greedy algorithm.


- Find the first non-empty bin.
- Move earth into first non-empty bin in the other histogram.



Earth Mover's Distance - illustration

- In general, this is the *transportation problem* and can be solved using linear programming.
- For 1D histograms, this can be solved using the greedy algorithm.


- Find the first non-empty bin.
- Move earth into first non-empty bin in the other histogram.



Earth Mover's Distance - illustration

- In general, this is the *transportation problem* and can be solved using linear programming.
- For 1D histograms, this can be solved using the greedy algorithm.


- Find the first non-empty bin.
- Move earth into first non-empty bin in the other histogram.



Earth Mover's Distance - illustration

- In general, this is the *transportation problem* and can be solved using linear programming.
- For 1D histograms, this can be solved using the greedy algorithm.


- Find the first non-empty bin.
- Move earth into first non-empty bin in the other histogram.



Earth Mover's Distance - illustration

- In general, this is the *transportation problem* and can be solved using linear programming.
- For 1D histograms, this can be solved using the greedy algorithm.


- Find the first non-empty bin.
- Move earth into first non-empty bin in the other histogram.



Earth Mover's Distance - illustration

- In general, this is the *transportation problem* and can be solved using linear programming.
- For 1D histograms, this can be solved using the greedy algorithm.


- Find the first non-empty bin.
- Move earth into first non-empty bin in the other histogram.



Earth Mover's Distance - illustration

- In general, this is the *transportation problem* and can be solved using linear programming.
- For 1D histograms, this can be solved using the greedy algorithm.


- Find the first non-empty bin.
- Move earth into first non-empty bin in the other histogram.



Earth Mover's Distance - illustration

- In general, this is the *transportation problem* and can be solved using linear programming.
- For 1D histograms, this can be solved using the greedy algorithm.

- Find the first non-empty bin.
- Move earth into first non-empty bin in the other histogram.



Earth Mover's Distance

Alternatively:

Compute the cumulative distributions:

$$CDX(i) = \sum_{j=1}^i x_j$$

Then the Earth Mover's Distance between X and Y is:

$$EMD(X, Y) = \sum_{i=1}^n |CDX(i) - CDY(i)|$$

So that for 1D histograms, the EMD can be expressed as a normed difference.

Distance/Similarity measures

- Euclidean L_2
- Manhattan L_1
- Chessboard L_{inf}
- Euclidean L_2
- Cosine
- Correlation
- Intersection
- Kullback-Leibler divergence
- Jenses-Shanon divergence
- Chi square
- Quadratic form distance
- Earth movers distance
- **Mahalanobis distance**

Distance Metric Learning

Distance Metric Learning - motivation

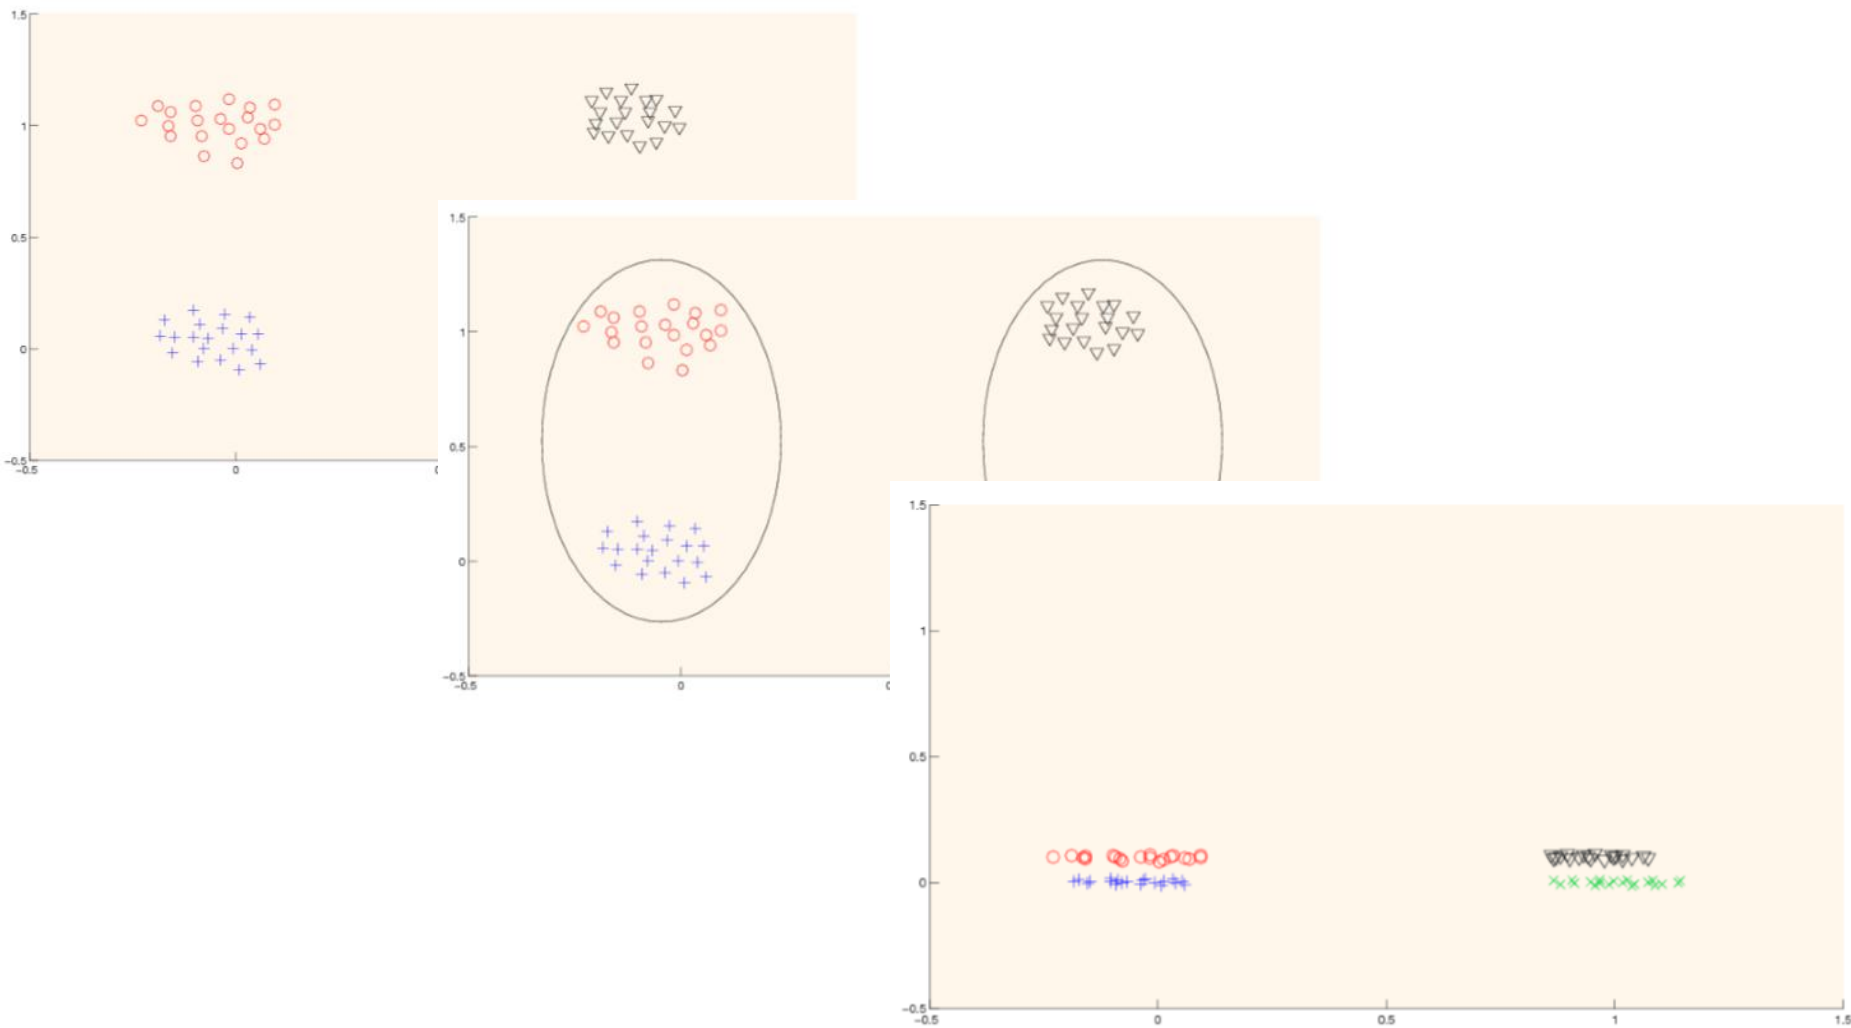
Learn metric for scaling dimensions/features

- Example: UCI Wine data set
 - 13 features 9/13 features have mean value in $[0,10]$
 - 3/13 features have mean value in $[10,100]$
 - One feature has a mean value of 747 (with std 315)
- Using a standard distance such as Euclidean distance, the largest feature dominates the computation
 - That feature may not be important for classification
- Need a weighting of the features that improves classification or other tasks

- 1) Alcohol
- 2) Malic acid
- 3) Ash
- 4) Alcalinity of ash
- 5) Magnesium
- 6) Total phenols
- 7) Flavanoids
- 8) Nonflavanoid phenols
- 9) Proanthocyanins
- 10) Color intensity
- 11) Hue
- 12) OD280/OD315 of diluted wines
- 13) Proline

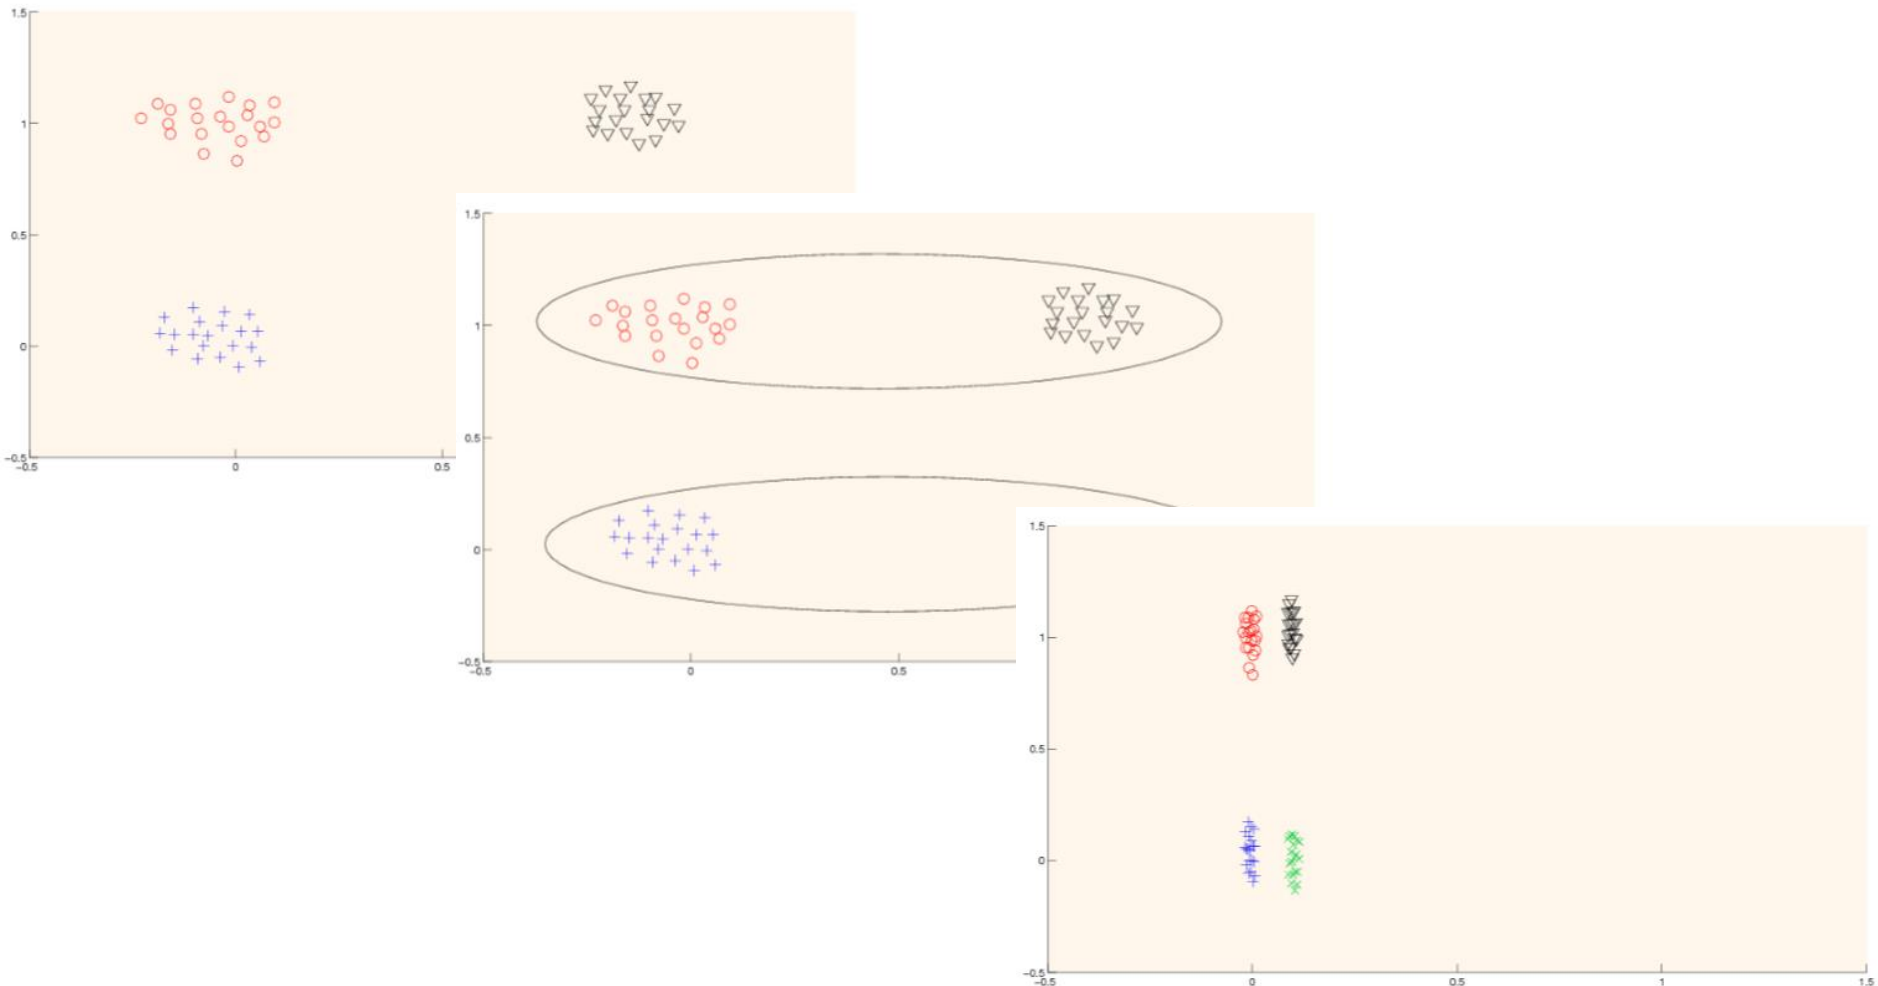
Distance Metric Learning - motivation

- Learn metric for reducing/grouping dimensions/features



Distance Metric Learning - motivation

- Learn metric for reducing/grouping dimensions/features



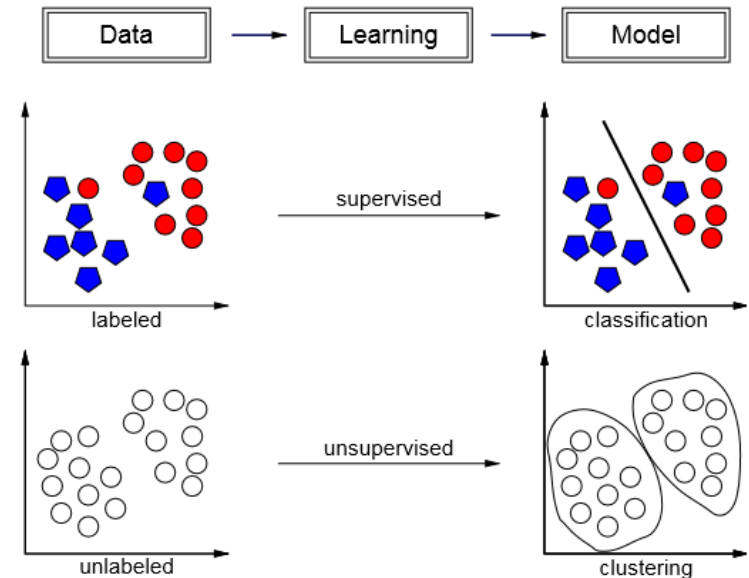
Distance Metric Learning

- Feature re-weighting/scaling
 - Learn weightings (scaling) over the features, then use standard distance (e.g., Euclidean) after re-weighting
 - Diagonal Mahalanobis methods
 - Number of parameters grows linearly with the dimensionality d
- Full linear transformation
 - In addition to scaling of features, also rotates the data
 - Linear dimensionality reduction for transformations to $r < d$ dimensions
 - For transformations from d dimensions to d dimensions, number of parameters grows quadratically in d
- Non-linear transformation
 - Variety of methods
 - Kernelization of linear transformations
 - Neural nets
 - Complexity varies from method to method

Distance Metric Learning

- Unsupervised Metric Learning

- Dimensionality reduction
- Principal Components Analysis
- Kernel PCA
- Multidimensional Scaling



- Supervised and Semi-supervised Metric Learning

- Constraints or labels given to the algorithm
- Example: set of similarity and dissimilarity constraints

Mahalanobis distance – original covariance based

ON THE GENERALIZED DISTANCE IN STATISTICS.

By P. C. MAHALANOBIS.

(Read January 4, 1936.)

1. A normal (Gauss-Laplacian) statistical population in P -variates is usually described by a P -dimensional frequency distribution :—

$$df = \text{const.} \times e^{-\frac{1}{2\alpha} \left[A_{11}(x_1 - \alpha_1)^2 + A_{22}(x_2 - \alpha_2)^2 + \dots + 2A_{12}(x_1 - \alpha_1)(x_2 - \alpha_2) + \dots \right]} \cdot dx_1 \cdot dx_2 \dots dx_P \quad (1.0)$$

where

$\alpha_1, \alpha_2, \dots, \alpha_P$ = the population (mean) values
of the P -variates x_1, x_2, \dots, x_P (1.1)

$\alpha_{ii} = \sigma_i^2$, are the respective variances (1.2)

$\alpha_{ij} = \sigma_i \cdot \sigma_j \cdot \rho_{ij}$, where ρ_{ij} = the coefficient of correlation between the i th and j th variates (1.3)

α is the determinant $|\alpha_{ij}|$ defined more fully in (2.2), and A_{ij} is the minor of α_{ij} in this determinant.

A P -variate normal population is thus completely specified by the set of $P(P+1)/2$ parameters * :—

Mahalanobis distance – original covariance based

Assume the data is represented as N vectors of length d :

$$X = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N]$$


Squared Euclidean distance

$$\begin{aligned} d(\mathbf{x}_1, \mathbf{x}_2) &= \|\mathbf{x}_1 - \mathbf{x}_2\|_2^2 \\ &= (\mathbf{x}_1 - \mathbf{x}_2)^T (\mathbf{x}_1 - \mathbf{x}_2) \end{aligned}$$

$$\text{Let } \Sigma = \sum_{i,j} (\mathbf{x}_i - \mu)(\mathbf{x}_j - \mu)^T$$

The “original” Mahalanobis distance:

$$d_M(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 - \mathbf{x}_2)^T \Sigma^{-1} (\mathbf{x}_1 - \mathbf{x}_2)$$



Mahalanobis distance

- Weighting/scaling of dimensions/features
 - Diagonal covariance matrix

$$f_1 = [v_1, \dots, v_N] \quad D_{Euclidean}^2(f_1, f_2) = \|f_1 - f_2\|^2 = \sum_n (f_1(n) - f_2(n))^2 \quad f_2 = [v_1, \dots, v_N]$$


$$D_{Mahalanobis}^2(f_1, f_2) = [f_1 - f_2]^T \cdot \Sigma^{-1} \cdot [f_1 - f_2] =$$

$$= \begin{bmatrix} f_1(0) - f_2(0) \\ f_1(1) - f_2(1) \\ \vdots \end{bmatrix}^T \begin{bmatrix} \sigma_1^2 & \sigma_{12} & \cdots \\ \sigma_{12} & \sigma_2^2 & \cdots \\ \vdots & \vdots & \ddots \end{bmatrix}^{-1} \begin{bmatrix} f_1(0) - f_2(0) \\ f_1(1) - f_2(1) \\ \vdots \end{bmatrix}$$

σ_a^2 – variance of component $f(a)$
 σ_{ab} – covariance between features $f(a), f(b)$

$$\text{if } \sigma_{ab} = 0 \quad \forall a \neq b$$

$$\text{then } D_{Mahalanobis}^2(f_1, f_2) = \sum_n \frac{(f_1(n) - f_2(n))^2}{\sigma_n^2}$$



$$\text{if } \sigma_{ab} = \text{const} \quad \forall a = b, \text{ and } \sigma_{ab} = 0 \quad \forall a \neq b$$

$$D_{Mahalanobis}^2(f_1, f_2) = \frac{1}{\text{const}^2} \sum_n (f_1(n) - f_2(n))^2 = D_{Euclidean}^2(f_1, f_2)$$

Mahalanobis distance - general

- General form

$$d_A(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 - \mathbf{x}_2)^T A (\mathbf{x}_1 - \mathbf{x}_2)$$

- Why is A positive semi-definite (PSD)?
 - If A is not PSD, then d_A could be negative
 - Suppose $\mathbf{v} = \mathbf{x}_1 - \mathbf{x}_2$ is an eigenvector corresponding to a negative eigenvalue λ of A

$$\begin{aligned} d_A(\mathbf{x}_1, \mathbf{x}_2) &= (\mathbf{x}_1 - \mathbf{x}_2)^T A (\mathbf{x}_1 - \mathbf{x}_2) \\ &= \mathbf{v}^T A \mathbf{v} \\ &= \lambda \mathbf{v}^T \mathbf{v} = \lambda < 0 \end{aligned}$$

Mahalanobis distance - general

- Properties of a metric:
 - $d(\mathbf{x}, \mathbf{y}) \geq 0$
 - $d(\mathbf{x}, \mathbf{y}) = 0$ if and only if $\mathbf{x} = \mathbf{y}$
 - $d(\mathbf{x}, \mathbf{y}) = d(\mathbf{y}, \mathbf{x})$
 - $d(\mathbf{x}, \mathbf{z}) \leq d(\mathbf{x}, \mathbf{y}) + d(\mathbf{y}, \mathbf{z})$
- d_A is not technically a metric $d_A(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 - \mathbf{x}_2)^T A (\mathbf{x}_1 - \mathbf{x}_2)$
 - Analogous to Euclidean distance, need the square root:


$$\sqrt{d_A(\mathbf{x}_1, \mathbf{x}_2)} = \sqrt{(\mathbf{x}_1 - \mathbf{x}_2)^T A (\mathbf{x}_1 - \mathbf{x}_2)}$$

- Square root of the Mahalanobis distance satisfies all properties if A is strictly positive definite, but if A is positive semi-definite then second property is not satisfied
 - Called a *pseudo-metric*
- In practice, most algorithms work only with d_A

Mahalanobis distance - general

- Can view d_A as the squared Euclidean distance after applying a linear transformation
 - Decompose $A = G^T G$ via Cholesky decomposition
 - (Alternatively, take eigenvector decomposition $A = V\Lambda V^T$ and look at $A = (\Lambda^{1/2} V^T)^T (\Lambda^{1/2} V^T)$)

$$\begin{aligned}d_A(\mathbf{x}_1, \mathbf{x}_2) &= (\mathbf{x}_1 - \mathbf{x}_2)^T A (\mathbf{x}_1 - \mathbf{x}_2) \\&= (\mathbf{x}_1 - \mathbf{x}_2)^T G^T G (\mathbf{x}_1 - \mathbf{x}_2) \\&= (G\mathbf{x}_1 - G\mathbf{x}_2)^T (G\mathbf{x}_1 - G\mathbf{x}_2) \\&= \|G\mathbf{x}_1 - G\mathbf{x}_2\|_2^2\end{aligned}$$



- Mahalanobis distance is just the squared Euclidean distance after applying the linear transformation G
 - Unitless and scale invariant



Mahalanobis distance learning - illustration

- Given this grouping we may want to learn matrix A such that it brings closer the related groups

$$\begin{aligned}d_A(\mathbf{x}_1, \mathbf{x}_2) &= (\mathbf{x}_1 - \mathbf{x}_2)^T A (\mathbf{x}_1 - \mathbf{x}_2) \\&= (\mathbf{x}_1 - \mathbf{x}_2)^T G^T G (\mathbf{x}_1 - \mathbf{x}_2) \\&= (G\mathbf{x}_1 - G\mathbf{x}_2)^T (G\mathbf{x}_1 - G\mathbf{x}_2) \\&= \|G\mathbf{x}_1 - G\mathbf{x}_2\|_2^2\end{aligned}$$

learn:

$$A = \begin{pmatrix} 1 & 0 \\ 0 & \epsilon \end{pmatrix} \quad G = \begin{pmatrix} 1 & 0 \\ 0 & \sqrt{\epsilon} \end{pmatrix}$$





Mahalanobis distance learning - illustration

- Given this grouping we may want to learn matrix A such that it brings closer the related groups

$$\begin{aligned}d_A(\mathbf{x}_1, \mathbf{x}_2) &= (\mathbf{x}_1 - \mathbf{x}_2)^T A (\mathbf{x}_1 - \mathbf{x}_2) \\&= (\mathbf{x}_1 - \mathbf{x}_2)^T G^T G (\mathbf{x}_1 - \mathbf{x}_2) \\&= (G\mathbf{x}_1 - G\mathbf{x}_2)^T (G\mathbf{x}_1 - G\mathbf{x}_2) \\&= \|G\mathbf{x}_1 - G\mathbf{x}_2\|_2^2\end{aligned}$$

learn:

$$A = \begin{pmatrix} \epsilon & 0 \\ 0 & 1 \end{pmatrix} \quad G = \begin{pmatrix} \sqrt{\epsilon} & 0 \\ 0 & 1 \end{pmatrix}$$



Mahalanobis distance learning - formulation

- Typically 2 main pieces to a Mahalanobis metric learning problem
 - A set of constraints on the distance
 - A regularizer on the distance / objective function
- In the constrained case, a general problem may look like:

$$\begin{array}{ll} \min_A & r(A) \\ \text{s.t.} & c_i(A) \leq 0 \quad 0 \leq i \leq C \\ & A \succeq 0 \end{array}$$

- r is a regularizer/objective on A and c_i are the constraints on A

Mahalanobis distance learning - constraints

- Similarity / Dissimilarity constraints
 - Given a set of pairs \mathcal{S} of points that should be similar, and a set of pairs of points \mathcal{D} of points that should be dissimilar
 - A single constraint would be of the form

$$d_A(\mathbf{x}_i, \mathbf{x}_j) \leq \ell$$

for $(i, j) \in \mathcal{S}$ or


$$d_A(\mathbf{x}_i, \mathbf{x}_j) \geq u$$

for $(i, j) \in \mathcal{D}$

- Easy to specify given class labels
- Aggregate distance constraints
 - Constrain the sum of all pairs of same-class distances to be small, e.g.,

$$\sum_{ij} y_{ij} d_A(\mathbf{x}_i, \mathbf{x}_j) \leq 1$$

where $y_{ij} = 1$ if \mathbf{x}_i and \mathbf{x}_j are in the same class, and 0 otherwise



Mahalanobis distance learning - constraints

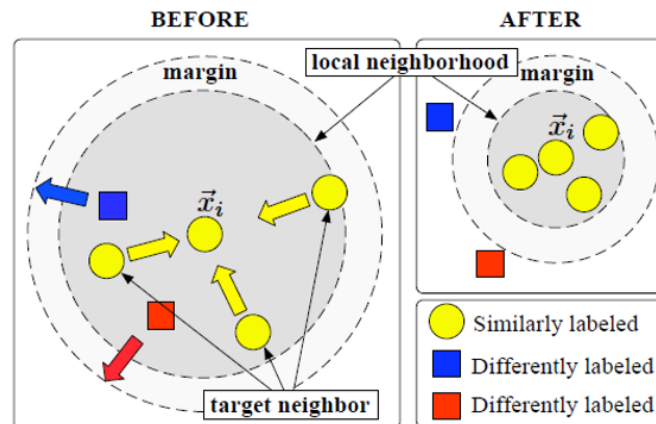
- Relative distance constraints

- Given a triple $(\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k)$ such that the distance between \mathbf{x}_i and \mathbf{x}_j should be smaller than the distance between \mathbf{x}_i and \mathbf{x}_k , a single constraint is of the form

$$d_A(\mathbf{x}_i, \mathbf{x}_j) \leq d_A(\mathbf{x}_i, \mathbf{x}_k) - m,$$

where m is the margin

- Popular for ranking problems



Define constraints tailored to k-NN in a local way

- the k nearest neighbours should be of same class ("target neighbors"), while examples of different classes should be kept away ("impostors"):

$$S = \{(x_i, x_j) : y_{ij} = 1 \text{ and } x_j \text{ belongs to the } k\text{-neighborhood of } x_i\},$$

$$R = \{(x_i, x_j, x_k) : y_{ij} = 1, y_{ik} = 0\}.$$

Mahalanobis distance learning - objectives

- General loss functions $D(A, A_0)$
 - Distance between matrices A and A_0
 - A is a matrix we learn
 - A_0 is a baseline matrix, which we start the learning with
- Loss/divergence functions
 - Squared Frobenius $\|A - A_0\|_F^2$
 - LogDet divergence: $\text{tr}(AA_0^{-1}) - \log \det(AA_0^{-1}) - d$
 - $\|A\|_F^2$ Frobenius norm
 - $\text{tr}(AC_0)$ (i.e., if C_0 is the identity, this is the trace norm)

Mahalanobis distance learning - objectives

- Frobenius norm – “Euclidean” norm used for matrices

$$\|A\|_F \equiv \sqrt{\sum_{i=1}^m \sum_{j=1}^n |a_{ij}|^2}$$

- LogDet divergence

$$\begin{aligned} D_{\ell d}(X, Y) &= \text{trace}(XY^{-1}) - \log \det(XY^{-1}) - d, \\ &= \text{trace}(Y^{-1/2}XY^{-1/2}) - \log \det(Y^{-1/2}XY^{-1/2}) - d. \end{aligned}$$

- Can be used as a measure of distance
 - Positive, and zero iff $X = Y$
 - But not symmetric, and triangle inequality does not hold
 - Scale-invariance

$$D_{\ell d}(X, Y) = D_{\ell d}(\alpha X, \alpha Y), \quad \alpha \geq 0$$

- Trace - sum of diagonal elements

$$\text{tr}(A) = \sum_i a_{ii}$$

Mahalanobis distance learning - objectives

- Example 1: $\text{tr}(A)$
 - Trace function is the sum of the eigenvalues
 - Analogous to the ℓ_1 penalty, promotes sparsity
 - Leads to low-rank A
- Example 2: LogDet Divergence
 - Defined only over positive semi-definite matrices
 - Makes computation simpler
 - Possesses other desirable properties
- Example 3: $\|A\|_F^2$
 - Arises in many formulations
 - Easy to analyze and optimize

Mahalanobis distance learning – example A

Problem posed as follows:

$$\begin{aligned} \max_A \quad & \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{D}} \sqrt{d_A(\mathbf{x}_i, \mathbf{x}_j)} \\ \text{s.t.} \quad & c(A) = \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{S}} d_A(\mathbf{x}_i, \mathbf{x}_j) \leq 1 \\ & A \succeq 0. \end{aligned}$$

- Here, \mathcal{D} is a set of pairs of dissimilar pairs, \mathcal{S} is a set of similar pairs
- Objective tries to maximize sum of dissimilar distances
- Constraint keeps **sum** of similar distances small
 - Use square root in regularizer to avoid trivial solution

[Xing, Ng, Jordan, and Russell; NIPS 2002]

Mahalanobis distance learning – example B

Problem formulated as follows:

$$\begin{aligned} \min_A \quad & \|A\|_F^2 \\ \text{s.t.} \quad & d_A(\mathbf{x}_i, \mathbf{x}_k) - d_A(\mathbf{x}_i, \mathbf{x}_j) \geq 1 \quad \forall (i, j, k) \in \mathcal{R} \\ & A \succeq 0. \end{aligned}$$

- Constraints in \mathcal{R} are relative distance constraints
- There may be no solution to this problem; introduce slack variables ξ_{ijk}

$$\begin{aligned} \min_{A, \xi} \quad & \|A\|_F^2 + \gamma \sum_{(i, j, k) \in \mathcal{R}} \xi_{ijk} \\ \text{s.t.} \quad & d_A(\mathbf{x}_i, \mathbf{x}_k) - d_A(\mathbf{x}_i, \mathbf{x}_j) \geq 1 - \xi_{ijk} \quad \forall (i, j, k) \in \mathcal{R} \\ & \xi_{ijk} \geq 0 \quad \forall (i, j, k) \in \mathcal{R} \\ & A \succeq 0. \end{aligned}$$

[Schultz and Joachims; NIPS 2002]

Mahalanobis distance learning – example C

- Problem Formulation

$$\begin{aligned} \min_A \quad & \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{S}} d_A(\mathbf{x}_i, \mathbf{x}_j) \\ \text{s.t.} \quad & d_A(\mathbf{x}_i, \mathbf{x}_k) - d_A(\mathbf{x}_i, \mathbf{x}_j) \geq 1 \quad \forall (\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R} \\ & A \succeq 0. \end{aligned}$$

- Also define set \mathcal{S} of pairs of points $(\mathbf{x}_i, \mathbf{x}_j)$ such that \mathbf{x}_i and \mathbf{x}_j are neighbors in the same class
- Want to minimize sum of distances of pairs of points in \mathcal{S}
- Also want to satisfy the relative distance constraints
- Introduce slack variables


$$\begin{aligned} \min_{A, \xi} \quad & \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{S}} d_A(\mathbf{x}_i, \mathbf{x}_j) + \gamma \sum_{(\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R}} \xi_{ijk} \\ \text{s.t.} \quad & d_A(\mathbf{x}_i, \mathbf{x}_k) - d_A(\mathbf{x}_i, \mathbf{x}_j) \geq 1 - \xi_{ijk} \quad \forall (\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R} \\ & A \succeq 0, \xi_{ijl} \geq 0. \end{aligned}$$

Mahalanobis distance learning

- Many existing Mahalanobis distance learning methods can be obtained simply by choosing a regularizer/objective and constraints
 - Memory overhead grows quadratically with the dimensionality of the data
 - Does not scale to high-dimensional data ($d = O(10^6)$ for many image embeddings)
 - Only works for linearly separable data

but

There is no linear transformation for such grouping problem





Kernel Distance learning - nonlinear

- Map input data to higher-dimensional “feature” space:

$$\mathbf{x} \rightarrow \varphi(\mathbf{x})$$

- Idea: Run machine learning algorithm in feature space
- Use the following mapping:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \rightarrow \begin{bmatrix} x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix}$$



Kernel Distance learning - nonlinear

- Map input data to higher-dimensional “feature” space:

$$\mathbf{x} \rightarrow \varphi(\mathbf{x})$$

- Idea: Run machine learning algorithm in feature space
- Use the following mapping:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \rightarrow \begin{bmatrix} x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix}$$

- Kernel function: $\kappa(\mathbf{x}, \mathbf{y}) = \langle \varphi(\mathbf{x}), \varphi(\mathbf{y}) \rangle$
- “Kernel trick” — no need to explicitly form high-dimensional features
- In this example: $\langle \varphi(\mathbf{x}), \varphi(\mathbf{y}) \rangle = (\mathbf{x}^T \mathbf{y})^2$

Kernel Distance learning - nonlinear

- Kernel function (Definition)
 - A symmetric similarity function K is a kernel if there exist a (possibly implicit) mapping function $\phi : X \rightarrow H$ from instance space X to a Hilbert space H such that K can be written as an inner product in H :
$$K(x_i, x_j) = \langle \phi(x_i), \phi(x_j) \rangle$$
 - K is a kernel **if** it is *positive semi-definite* i.e. the following

$$\sum_{i,j=1}^n c_i c_j K(x_i, x_j) \geq 0$$

holds for any $n \in \mathbb{N}$, $x_1, \dots, x_n \in \mathcal{X}$, $c_1, \dots, c_n \in \mathbb{R}$

Kernel Distance learning - nonlinear

- Main idea
 - Take an existing learning algorithm
 - Write it using inner products
 - Replace inner products $\mathbf{x}^T \mathbf{y}$ with kernel functions $\varphi(\mathbf{x})^T \varphi(\mathbf{y})$
 - If $\varphi(\mathbf{x})$ is a non-linear function, then algorithm has been *implicitly* non-linearly mapped
- Examples of kernel functions

$$\kappa(\mathbf{x}, \mathbf{y}) = (\mathbf{x}^T \mathbf{y})^p \quad \text{Polynomial Kernel}$$

$$\kappa(\mathbf{x}, \mathbf{y}) = \exp\left(-\frac{\|\mathbf{x} - \mathbf{y}\|_2^2}{2\sigma^2}\right) \quad \text{Gaussian Kernel}$$

$$\kappa(\mathbf{x}, \mathbf{y}) = \tanh(c(\mathbf{x}^T \mathbf{y}) + \theta) \quad \text{Sigmoid Kernel}$$

- Kernel functions also defined over objects such as images, trees, graphs, etc.

Kernel Distance learning - kmeans example

Recall the k -means clustering algorithm

- Repeat until convergence:
 - Compute the means of every cluster π_c

$$\mu_c = \frac{1}{|\pi_c|} \sum_{\mathbf{x}_i \in \pi_c} \mathbf{x}_i$$

- Reassign points to their closest mean by computing

$$\|\mathbf{x} - \mu_c\|_2^2$$

for every data point \mathbf{x} and every cluster π_c

Kernelization of k -means

- Expand $\|\mathbf{x} - \mu_c\|_2^2$ as

$$\mathbf{x}^T \mathbf{x} - \frac{2 \sum_{\mathbf{x}_i \in \pi_c} \mathbf{x}^T \mathbf{x}_i}{|\pi_c|} + \frac{\sum_{\mathbf{x}_i, \mathbf{x}_j \in \pi_c} \mathbf{x}_i^T \mathbf{x}_j}{|\pi_c|^2}$$
 - No need to explicitly compute the mean; just compute this for every point to every cluster
- Replace inner products with kernels, and this is kernel k -means

$$\kappa(\mathbf{x}, \mathbf{x}) - \frac{2 \sum_{\mathbf{x}_i \in \pi_c} \kappa(\mathbf{x}, \mathbf{x}_i)}{|\pi_c|} + \frac{\sum_{\mathbf{x}_i, \mathbf{x}_j \in \pi_c} \kappa(\mathbf{x}_i, \mathbf{x}_j)}{|\pi_c|^2}$$

Kmeans finds linear separators for cluster boundaries,
kernel Kmeans finds non-linear separators

Distance learning – Mahalanobis kernel summary

Consider the following *kernelized* problem

- You are given a kernel function $\kappa(\mathbf{x}, \mathbf{y}) = \varphi(\mathbf{x})^T \varphi(\mathbf{y})$
- You want to run a metric learning algorithm in kernel space
 - Optimization algorithm cannot use the explicit feature vectors $\varphi(\mathbf{x})$
 - Must be able to compute the distance/kernel over arbitrary points (not just training points)

- Mahalanobis distances:

$$d_A(\mathbf{x}, \mathbf{y}) = (\mathbf{x} - \mathbf{y})^T A (\mathbf{x} - \mathbf{y})$$

$$d_A(\mathbf{x}, \mathbf{y}) = (\varphi(\mathbf{x}) - \varphi(\mathbf{y}))^T A (\varphi(\mathbf{x}) - \varphi(\mathbf{y}))$$

- Inner products / kernels:

$$\kappa_A(\mathbf{x}, \mathbf{y}) = \mathbf{x}^T A \mathbf{y}$$

$$\kappa_A(\mathbf{x}, \mathbf{y}) = \varphi(\mathbf{x})^T A \varphi(\mathbf{y})$$

- Algorithms for constructing A learn both measures

- Can be thought of as a kind of *kernel learning* problem